



University
of Regina

CS 215

Web Oriented Programming

Database Fundamentals

Dr. Orland Hoeber

orland.hoeber@uregina.ca

<http://www.cs.uregina.ca/~hoeber/cs215/>

Persistent Data Storage For Web Apps

- If we want to maintain data for any non-trivial length of time, we need to store it on the server
- Compared to storing it on the client (browser), storing it on the server is more:
 - ▣ reliable
 - ▣ maintainable
 - ▣ secure
- In almost all web application cases, such data is stored in a database
 - ▣ the goal of this lecture is to explain to you why this is the case

The Problems with File Storage

- Suppose we have a collection of data that is relatively large (e.g., 500 GB of customer records)
- If we want to access some of this within a web application, we will need to manage a few details:
 - ▣ load the entire dataset into memory
 - ▣ If the data is too large, devise a scheme for only loading relevant elements into memory
 - ▣ develop a scheme for addressing the data elements
 - ▣ write software that will support adding, editing, or changing the data
 - ▣ ensure that the software will support concurrent access and modifications
 - ▣ ensure that the software will handle crashes and reboots gracefully
- There is a lot of software development overhead here

Database Management Systems

- Database management systems (DBMS) take care of all of these problems for you
 - Data independence
 - data representation and storage are handled independently of the application program
 - Efficient data access
 - data storage and retrieval can occur with minimal delay
 - Data integrity and security
 - rules regarding valid data can be automatically enforced
 - access to the data can be restricted depending on who or what application is accessing it

DBMS (continued)

- ▣ Data administration
 - the data can be administered independent of the application
- ▣ Concurrent access
 - multiple users can access the data at the same time
 - maintains the illusion that each has dedicated access
- ▣ Crash recovery
 - if the computer crashes, no data is corrupted or lost
 - any changes to the data either happen completely or don't happen at all
- Fundamentally, the use of a DBMS will result in significant reductions in application development time
- However, in some special cases, databases may not be appropriate (e.g., real-time data processing)

Data Model

- A **data model** is a set of high-level data description constructs that allow us to specify what the data should look like without having to worry about the low-level storage details
- The specific data model used in most commonly is the **relational data model**
 - ▣ the central data construct in this model is the relation, which can be thought of as a set of records
 - ▣ don't confuse this with a relationship, which is also present in the relational model

Schema

- A specific description of data in terms of a data model is a **schema**
- Within a relational data model, the schema for a relation specifies:
 - ▣ the name of the relation
 - ▣ the name of each field (or attribute, or column)
 - ▣ the type of each field

```
Students (sid: integer, name: string, login: string,  
dob: date, gpa: float)
```

- Logically, this gives us a template for defining a table structure, where each row is a record that describes a student

Levels of Abstraction

- The data in a database can be described at three levels of abstraction:
 - ▣ conceptual
 - ▣ physical
 - ▣ external
- These all represent the database, but from different perspectives and purposes

Conceptual Schema

- The **conceptual schema** describes the database in terms of the data model
 - ▣ information about entities
 - ▣ relationships between the entities
- All of this information can be stored in a series of relations (tables)
- The choice of relations and fields in a relation is not always obvious
- The process of arriving at a good conceptual schema is called **conceptual database design**

Conceptual Database Design

Student (sid: integer, name: string, login: string,
dob: date, gpa: float)

Faculty (fid: integer, name: string)

Courses (cid: integer, cname: string,
credits:integer)

Enrolled (sid: integer, cid: integer, grade: float)

Teaches (fid: integer, cid: integer)

Physical Schema

- The **physical schema** specifies storage details for the conceptual schema
 - ▣ which file organizations to use to store the relations
 - ▣ what auxiliary data structures can be created to speed up storage and/or retrieval operations (indexes)
- Decisions on the physical schema are made based on an understanding of how the data is typically accessed

External Schema

- **External schemas** allow data access to be customized and authorized at the level of individual users and groups
 - ▣ different users can share access to the same database, but with different rights and privileges
 - ▣ a simple example is table-level access controls
 - ▣ a more complex example is custom views that combine different elements from different related tables
 - e.g., a view that shows student enrollment in a course, their grades, but not their login information
- Such external schemas are guided by the end-user requirements for the application

Data Independence

- A fundamental feature of using a DBMS within any application is that it offers **data independence**
 - ▣ application programs are insulated from changes in the way the data is structured and stored
- There are two different types of data independence:
 - ▣ logical data independence
 - users are shielded from changes in the logical structure of the data or changes in the relations
 - achieved through the use of views
 - ▣ physical data independence
 - changes in how the data is stored internally within the database have no impact on how the users access the data

Queries

- What makes databases popular is the ease by which information can be accessed
- Relational databases allow for a rich class of questions to be easily posed to the database
 - ▣ such questions are called **queries**
 - ▣ they are written in a **query language**
 - ▣ in addition, users can create, modify, and delete data through a **data manipulation language**
- The language we will use for these purpose is called the **structured query language**, or **SQL**

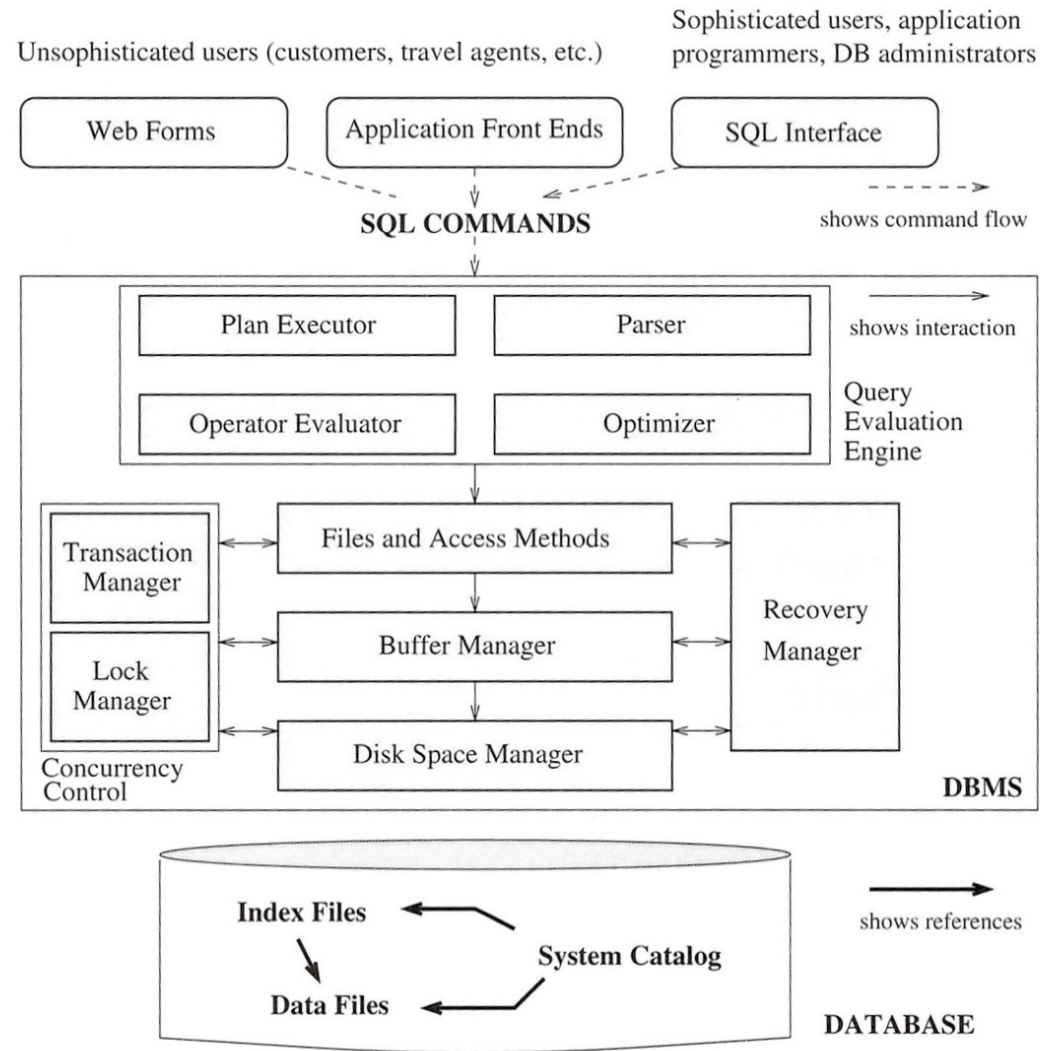
Transaction Management

- An important feature of databases is the ability to allow concurrent users of the database
- In order to avoid conflicts, the database must have a concept of the basic unit of change: the **transaction**
- The illusion of concurrent access is achieved by carefully interleaving the transactions of each user
- Problems are avoided by using a locking protocol on the relations
 - ▣ queries can use a **shared lock**
 - ▣ modifications require the use of an **exclusive lock**

Incomplete Transactions & Crashes

- When a transaction is interrupted (e.g., due to a crash), it is important to be able to recover from it gracefully
- This is done using a **write-ahead log**
 - ▣ write the details of the transaction to the log first
 - ▣ force the log to be written to the disk
 - ▣ execute the transaction on the database
- In this way, if a crash occurs, the DBMS can check the log and the status of the database to find out if the last few transactions were completed or not

Components & Structure



Databases & Tables (Relations)

- Tables

- ▣ data is organized over rows and columns

- Columns

- ▣ columns are named
- ▣ assigned specific data types

user_id	username	password	last_access
1	tim	xxkkddi	07-11-14
2	jill	33kk99d	05-11-14

- Rows

- ▣ rows are unnamed
- ▣ values are (normally) assigned for each column
- ▣ the rows represent entities, where the column values are the attributes for the entities

- Primary Keys

- ▣ most tables will have a special column that contains a unique identifier for each row – the primary key
- ▣ normally an integer value

Simple Database Tables

- The simplest database table will contain a primary key, and a column for each attribute of the data
 - ▣ this can result in a lot of wasted space and redundant data
 - ▣ a better approach is to think about the meaning of the data in a careful and logical manner, decomposing it into multiple related tables that eliminate the redundancy

car_id	make	model	engine	hp	doors	price
1	Chevrolet	Volt	electric	149	4	\$28,000
2	Audi	A4	4 cyl turbo	210	4	\$44,000

Normalization

- The process of separating your data into tables and creating keys that link them together is called normalization
- Goal: ensure that each piece of information appears in the database only once
 - ▣ duplicate data is a waste of space
 - ▣ when data that is duplicated and the database needs to be updated, there is a risk of the data becoming inconsistent (updated in one location, but not others)
- There are three standard forms of normalization: first, second, and third

First Normal Form (1NF)

- The database must fulfill the following requirements:
 - ▣ there should be no repeating columns containing the same data
 - ▣ all columns should contain a single value
 - ▣ there should be a primary key to uniquely identify each row
- This is simple to do with careful conceptual database design

Second Normal Form (2NF)

- To achieve second normal form, all the tables must start in first normal form
- Then, the goal is to eliminate redundancy across multiple rows by creating new tables that hold this repeating information once
 - e.g., suppose we have a database table that holds information about product orders. For repeat purchases, we should not duplicate the name and address of the customer within the orders table, but instead create a new table that holds customers independently from their orders.

Third Normal Form (3NF)

- Once you have the database in both first and second normal form, it is probably in good shape and may not need any further modifications
- Third normal form is the process of eliminating all other redundancy from the database
 - ▣ any data that is dependent on another value in the table should also be moved to a separate table
 - ▣ e.g., a city can be derived from the postal code, so this should be in a separate table
- **For our purposes, 3NF will result in too many tables and rather complex queries**

Entity Relationship Diagrams (ERD)

- ERDs are a useful tool for modeling the conceptual schema associated with a given problem or domain
 - ▣ describes the data (entities and attributes of the entities)
 - ▣ describe the relationships between the entities
- We can use language constructs to identify the features of an ERD
 - ▣ nouns: entities
 - ▣ adjectives: attributes
 - ▣ verbs: relationships




Entities

- Entities are the basic objects of ERDs
 - ▣ represent a “thing” with an independent existence
 - ▣ may be physical or conceptual
 - car, car type, manufacturer, engine
 - student, course, faculty, instructor
- Represented as blocks in the ERD



Student

Relationships

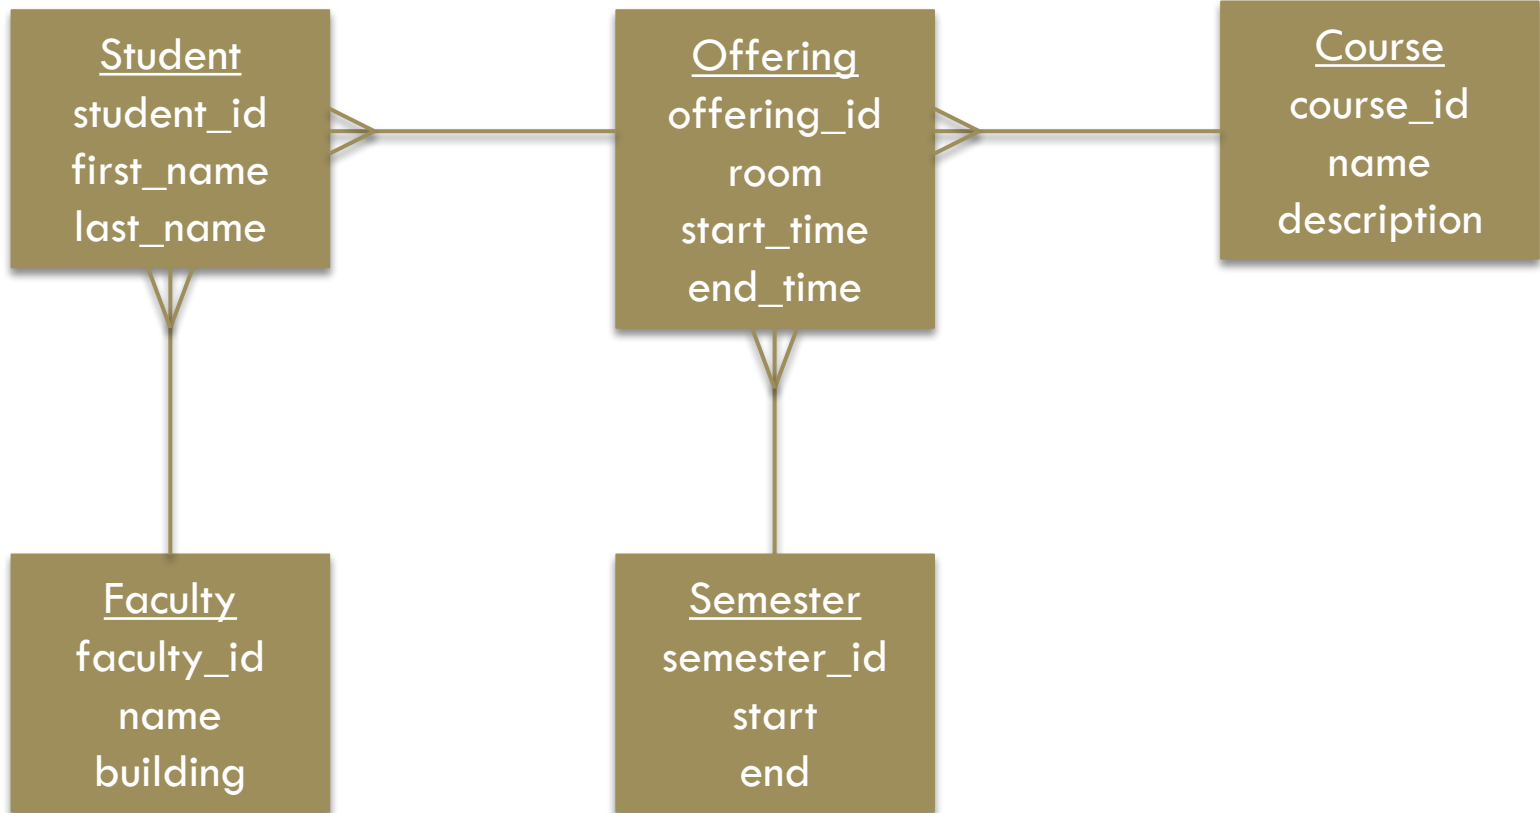
- Relationships are the associations between the entities
 - ▣ can be inclusion (like in OOO programming) or a more abstract connection
- There are three different types of relationships:
 - ▣ one to one: 
 - ▣ one to many: 
 - ▣ many to many: 

Attributes

- Attributes can be assigned to the entities and/or the relationships
 - ▣ for entity attributes, these can be listed within the entity box
 - ▣ relationship attributes are normally added as bubbles connected to the relationship line

A Simple Example

- Suppose we are modeling data for students attending courses across different faculties:
 - ▣ students can take multiple courses
 - ▣ courses are offered in semesters
 - ▣ students belong to faculties
- What are the entities?
- What are the attributes?
- What are the relationships?



The Value of ERDs

- The main value of carefully constructing an ERD is that it can readily be converted into a database structure:
 - ▣ entities become tables
 - ▣ relationships are represented with keys from one table into another
 - ▣ many to many relationships require separate tables to link the two entity tables

Final Words on Data Models

- Although our focus is on databases that implement the relational data model, there are other types of data models that are useful in special circumstances:
 - ▣ hierarchical data model
 - ▣ network data model
 - ▣ object-oriented data model
 - ▣ object-relational data model
 - ▣ semi-structured data model
- We have seen one of these data models in this class already (DOM), and will see another very soon (JSON)

Homework

- Read Chapter 8 & 9
- Next topic: Databases & MySQL
- Upcoming deadlines:
 - Assignment 4: Tuesday Nov 21 @ 11:55PM